

Project1

February 26, 2020

Start by importing the packages we need

```
[ ]: import pandas as pd
import numpy as np
import os
import sklearn
import matplotlib
```

Next step is to read in the data

```
[2]: myData = pd.read_csv("~/Documents/Titanic/pp-complete.csv", header = None)
```

Lets check out the data to understand it better

```
[3]: myData.head()
myData.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25018011 entries, 0 to 25018010
Data columns (total 16 columns):
#   Column  Dtype
---  -
0    0      object
1    1      int64
2    2      object
3    3      object
4    4      object
5    5      object
6    6      object
7    7      object
8    8      object
9    9      object
10   10     object
11   11     object
12   12     object
13   13     object
14   14     object
15   15     object
```

```
dtypes: int64(1), object(15)
memory usage: 3.0+ GB
```

Select variables of interest and rename them to make it more clear

```
[4]: modelData = myData[[1,2,4,6,11]]
modelData = modelData.rename(columns={1: "Price", 2: "Date", 4: "Type", 6: "Duration", 11: "Location"})
modelData.info()
modelData.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25018011 entries, 0 to 25018010
Data columns (total 5 columns):
#   Column      Dtype
---  -
0   Price       int64
1   Date        object
2   Type        object
3   Duration    object
4   Location    object
dtypes: int64(1), object(4)
memory usage: 954.4+ MB
```

```
[4]:
```

	Price	Date	Type	Duration	Location
0	18500	1995-01-31 00:00	F	L	TORQUAY
1	73450	1995-10-09 00:00	D	F	LIVERPOOL
2	59000	1995-03-31 00:00	D	F	POOLE
3	31000	1995-12-04 00:00	D	F	WOODBIDGE
4	95000	1995-09-22 00:00	D	F	LICHFIELD

Format the date column and get the year from the date which will be used to split the dataset later

```
[5]: modelData['Date'] = pd.to_datetime(modelData['Date'], format = '%Y-%m-%d')

modelData['Year'] = modelData['Date'].dt.year

modelData.info()
modelData.head(10)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25018011 entries, 0 to 25018010
Data columns (total 6 columns):
#   Column      Dtype
---  -
0   Price       int64
1   Date        datetime64[ns]
2   Type        object
3   Duration    object
```

```

4   Location    object
5   Year        int64
dtypes: datetime64[ns](1), int64(2), object(3)
memory usage: 1.1+ GB

```

```

[5]:
   Price      Date Type Duration      Location  Year
0   18500 1995-01-31   F         L      TORQUAY  1995
1   73450 1995-10-09   D         F    LIVERPOOL  1995
2   59000 1995-03-31   D         F        POOLE  1995
3   31000 1995-12-04   D         F   WOODBRIDGE  1995
4   95000 1995-09-22   D         F    LICHFIELD  1995
5   45450 1995-02-28   S         F CHESTERFIELD  1995
6   96000 1995-10-27   S         F        EPSOM  1995
7   30000 1995-11-28   S         F   WEDNESBURY  1995
8  425000 1995-03-31   D         F      COBHAM  1995
9   89995 1995-06-30   D         F   NORMANTON  1995

```

Add dummy variable “one-hot encode variable” if location of house is in London, use DataFrame from pandas to do it

```

[6]: from pandas import DataFrame
modelData['isLondon'] = modelData['Location'].apply(lambda x: 1 if x == 'LONDON' else 0)

modelData.info()

```

```

   Price      Date Type Duration      Location  Year  isLondon
0   18500 1995-01-31   F         L      TORQUAY  1995         0
1   73450 1995-10-09   D         F    LIVERPOOL  1995         0
2   59000 1995-03-31   D         F        POOLE  1995         0
3   31000 1995-12-04   D         F   WOODBRIDGE  1995         0
4   95000 1995-09-22   D         F    LICHFIELD  1995         0
...   ...   ...   ...   ...   ...   ...   ...
25018006  410854 2019-07-18   D         F      HORLEY  2019         0
25018007  610000 2019-08-08   D         F    CATERHAM  2019         0
25018008   42500 2019-07-22   0         F   GUILDFORD  2019         0
25018009  353500 2019-08-02   0         F    CHERTSEY  2019         0
25018010 1185000 2019-08-09   D         F      SUTTON  2019         0

```

```

[25018011 rows x 7 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25018011 entries, 0 to 25018010
Data columns (total 7 columns):
#   Column      Dtype
---  -
0   Price      int64
1   Date       datetime64[ns]
2   Type       object

```

```

3   Duration  object
4   Location  object
5   Year      int64
6   isLondon  int64
dtypes: datetime64[ns](1), int64(3), object(3)
memory usage: 1.3+ GB


```

Make a histogram of the price to see the distribution, notice it has a heavy tail on the left

```

[7]: import matplotlib.pyplot as plt
      %matplotlib inline
      modelData['Price'].hist(bins=30, range=(0, 2000000))
      plt.show()

```



Make some bar charts for the categorical variables to see how common each category is

```


[8]: #Get some bar charts of the categorical variables
      modelData['Type'].value_counts().plot(kind='bar')

```

```


[8]: <matplotlib.axes._subplots.AxesSubplot at 0x7f6e03ad5dd0>

```



```
[9]: modelData['Duration'].value_counts().plot(kind='bar')
```

```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x7f6e030f4f90>
```



Get dummy variables for each type of house and duration and add it to the dataset, rename them also to avoid confusion between variables

```
[10]: typeDummy = pd.get_dummies(modelData['Type'])
modelData = modelData.drop('Type',axis = 1)
modelData = modelData.join(typeDummy)

modelData = modelData.rename(columns={'D': "TypeD", 'F': "TypeF", 'O': "TypeO", 'S': "TypeS", 'T': "TypeT"})

[11]: durationDummy = pd.get_dummies(modelData['Duration'])
modelData = modelData.drop('Duration',axis = 1)
modelData = modelData.join(durationDummy)

modelData = modelData.rename(columns = {'F': 'DurationF', 'L': 'DurationL', 'U': 'DurationU'})
```

Make a copy of the dataset which we will later split to train and test sets

```
[12]: cleanedData = modelData.copy()
cleanedData.info()
```

```
[13]: <class 'pandas.core.frame.DataFrame'>
RangeIndex: 25018011 entries, 0 to 25018010
Data columns (total 13 columns):
 #   Column      Dtype
---  -
 0   Price      int64
 1   Date       datetime64[ns]
 2   Location   object
 3   Year       int64
 4   isLondon   int64
 5   TypeD      uint8
 6   TypeF      uint8
 7   TypeO      uint8
 8   TypeS      uint8
 9   TypeT      uint8
10   DurationF  uint8
11   DurationL  uint8
12   DurationU  uint8
dtypes: datetime64[ns](1), int64(3), object(1), uint8(8)
memory usage: 1.1+ GB
```

Delete variables we will not use and split the dataset into a training and test set. The test set includes variables which were sold in 2015 while the training set includes all the other properties.

```
[14]: del cleanedData['Date']
      del cleanedData['Location']
```

```
[15]: #Split to test and training set, test data is data in december, rest is
      →training data
      trainData = cleanedData[cleanedData.Year != 2015]
      trainData.info()
      testData = cleanedData[cleanedData.Year == 2015]
      testData.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24008337 entries, 0 to 25018010
Data columns (total 11 columns):
#   Column      Dtype
---  -
0   Price       int64
1   Year        int64
2   isLondon    int64
3   TypeD       uint8
4   TypeF       uint8
5   TypeO       uint8
6   TypeS       uint8
7   TypeT       uint8
8   DurationF   uint8
9   DurationL   uint8
10  DurationU   uint8
dtypes: int64(3), uint8(8)
memory usage: 915.8 MB
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1009674 entries, 20008866 to 21018539
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Price       1009674 non-null  int64
1   Year        1009674 non-null  int64
2   isLondon    1009674 non-null  int64
3   TypeD       1009674 non-null  uint8
4   TypeF       1009674 non-null  uint8
5   TypeO       1009674 non-null  uint8
6   TypeS       1009674 non-null  uint8
7   TypeT       1009674 non-null  uint8
8   DurationF   1009674 non-null  uint8
9   DurationL   1009674 non-null  uint8
10  DurationU   1009674 non-null  uint8
dtypes: int64(3), uint8(8)
memory usage: 38.5 MB
```

Delete variables we will not use

```
[16]: del trainData['Year']
      del testData['Year']
```

Make a vector for the variables for property price which is the variable we want to predict

```
[27]: trainPrice = trainData['Price']
      testPrice = testData['Price']
```

Fit the model, the model of choice is the random forest model.

```
[19]: from sklearn.ensemble import RandomForestRegressor
      ranForReg = RandomForestRegressor(n_estimators=10,n_jobs=-1)
      ranForReg.fit(trainData, trainData['Price'])
```

```
[19]: RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                             max_depth=None, max_features='auto', max_leaf_nodes=None,
                             max_samples=None, min_impurity_decrease=0.0,
                             min_impurity_split=None, min_samples_leaf=1,
                             min_samples_split=2, min_weight_fraction_leaf=0.0,
                             n_estimators=10, n_jobs=-1, oob_score=False,
                             random_state=None, verbose=0, warm_start=False)
```

Find the mean squared errors of the predictions compared to the actual observations.

```
[21]: from sklearn.metrics import mean_squared_error
      myPredictions = ranForReg.predict(trainData)
      ranForMSE = mean_squared_error(trainPrice, myPredictions)
      ranForRMSE = np.sqrt(ranForMSE)
```

```
[21]: 0.04071408274912438
```

```
[23]: ranForRMSE.round()
```

```
[23]: 7833.0
```

```
[31]: ranForRMSE/ trainPrice.mean()
```

```
[31]: 0.04071408274912438
```

```
[ ]: #Lets see how well the algorithm predicts by using 10-fold cross-validation
      import numpy as np
      from sklearn.model_selection import cross_val_score
      myScores = cross_val_score(ranForReg, trainData, trainPrice,
      ↪scoring="neg_mean_squared_error", cv=10,n_jobs=-1)
      ranForRmse = np.sqrt(-myScores)
```

```
[ ]:
```


[]:

[]:

[]:

[]:

[]:

```
[28]: #Get root squared mean errors for the test set, that is see how well the model  
      →predicts on the testing set and compare to the mean of the price  
finalPreds = ranForReg.predict(testData)  
finalMSE = mean_squared_error(testPrice, finalPreds)  
finalRMSE = np.sqrt(finalMSE)  
finalRMSE.round()
```

[28]: 7891.0

```
[29]: finalRMSE/testPrice.mean()
```

[29]: 0.02655437610770971

```
[32]: #Get root mean squared forecasting error  
SE = (finalPreds - testPrice) ** 2  
SFE = SE.divide(testPrice**2)  
MSFE = SFE.mean()  
RMSFE = np.sqrt(MSFE)  
RMSFE
```

[32]: 4.211414618902267e-05

[]: